


Experimental design and resource boundary

a

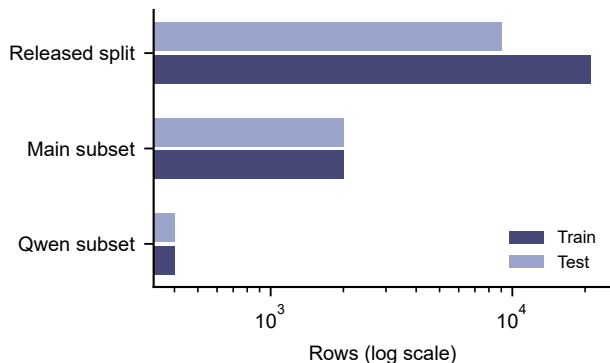
Scaled explanation-augmentation experiment on a single 8 GB GPU



No label or official explanation is passed to Qwen; test evaluation occurs once after training.

b

Data scales



c

Observed GPU envelope
Qwen: NVML total; RoBERTa: PyTorch allocated

